

Problems using deep generative models for probabilistic audio source separation

AMSTERDAM MACHINE LEARNING LAB
UNIVERSITY OF AMSTERDAM



Searching for PhD in
these topics!

Maurice Frank

@morris_frank__

maurice.frank@posteo.de







Maximilian Ilse

@MaxIlse

m.ilse@uva.nl

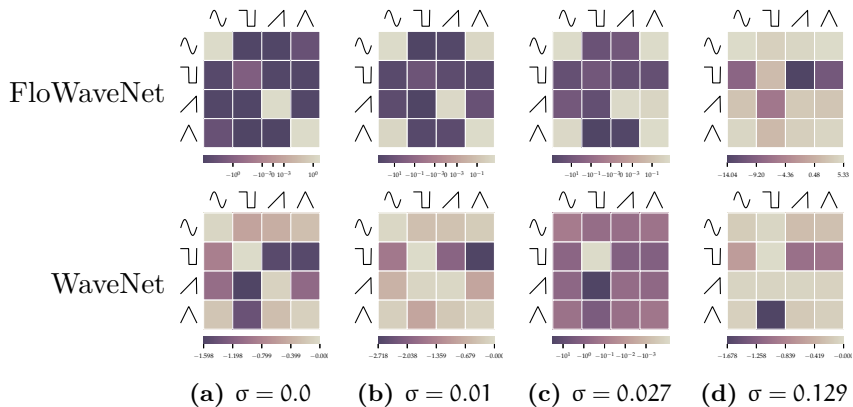
Cross-likelihood of the noiseless models



(a) Toy dataset

(b) *musdb18* dataset

Fine-tune with noisy samples



For your generative audio priors,
pick one:

Discriminative
but
peaked

or

Smooth
but
not-discriminative

